
General Agenda

1. Hardware interrupts
 2. Softirq
 3. Tasklet
 4. Workqueue
-


Detail Agenda

1. Interrupts
 2. Realtime and Interrupts
 3. Maskable interrupt
 4. Non-Maskable interrupt
 5. Exceptions
 6. APIC + PIC
 7. Softirq
 8. Tasklet
 9. Workqueue
 10. Detail
-

Mr. Micro-controller, May I have your attention?

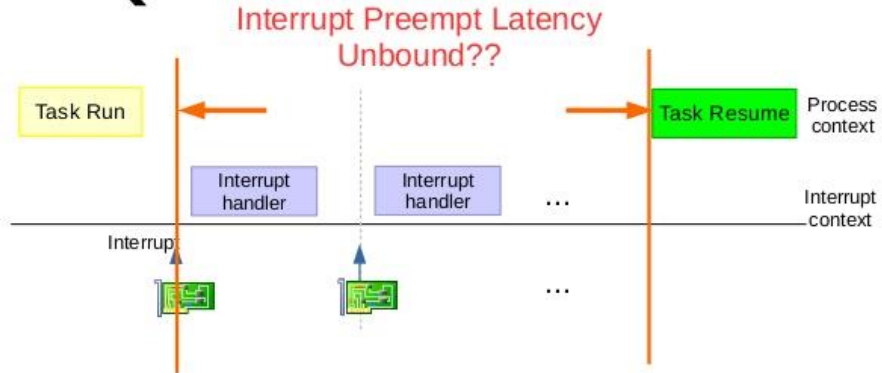


I am a peripheral!



Interrupts


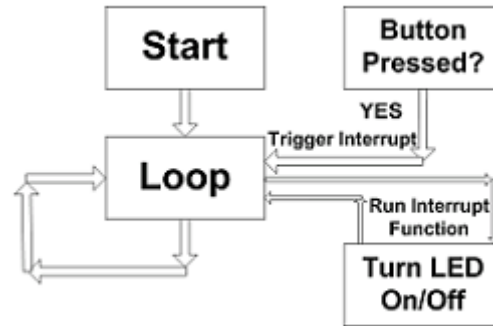
Task interrupted by IRQ




Interrupt

Polling

- **Difficult** is something that can be done **immediately**;
- **Impossible** is that which will take a little **longer**”






IRQ - Interrupt ReQuest

Interrupts


- **Difficult** is something that can be done **immediately**;
- **Impossible** is that which will take a little **longer**”





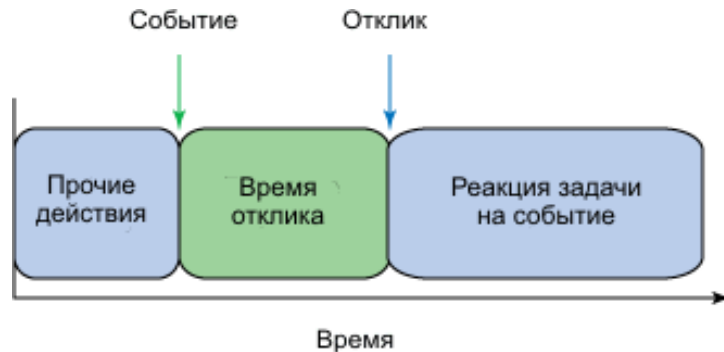
RealTime and Interrupts

- **Hard** – missing a deadline is a total system failure.
- **Firm** – infrequent deadline misses are tolerable, but may degrade the system's quality of service. The usefulness of a result is zero after its deadline.
- **Soft** – the usefulness of a result degrades after its deadline, thereby degrading the system's quality of service.



RealTime and Interrupts

- **Асинхронные события** — полностью непредсказуемые события. Например, вызов абонента телефонной станции.
- **Синхронные события** — предсказуемые события, случающиеся с определённой регулярностью. Например, вывод аудио и видео.
- **Изохронные события** — регулярные события (разновидность асинхронных), случающиеся в течение интервала времени. Например, в мультимедийном приложении данные аудиопотока должны прийти за время прихода соответствующей части потока видео.

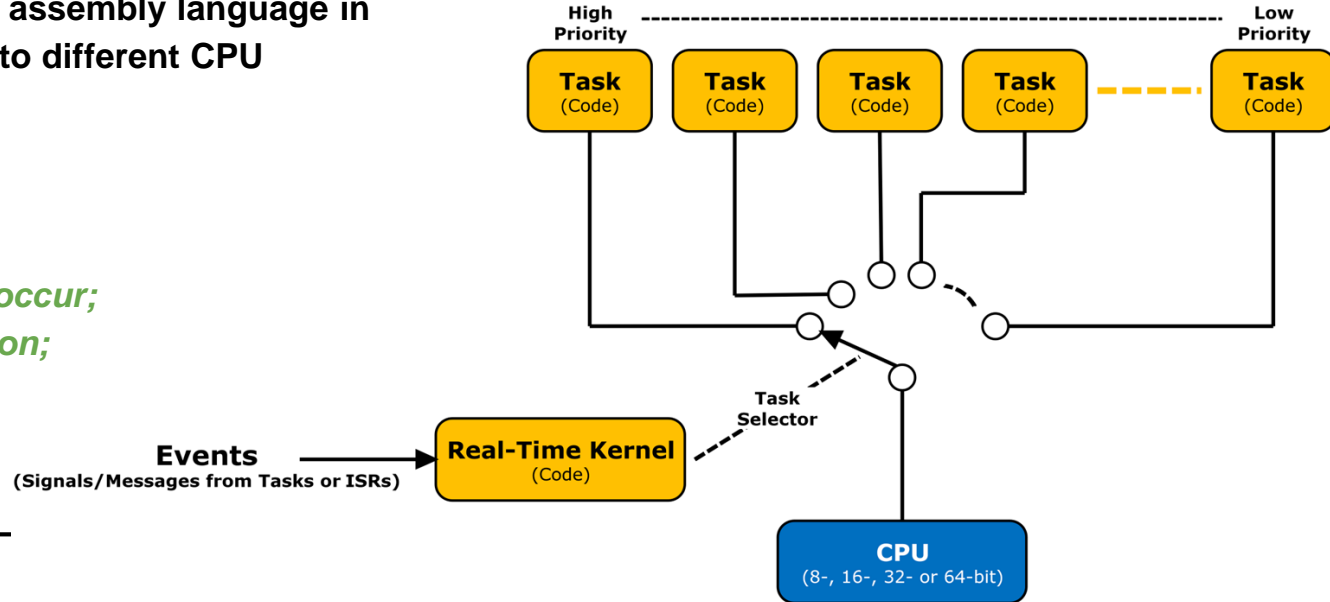


RealTime and Interrupts

- **Real-Time Kernel**
- A real-time kernel is software that manages the time of a CPU (Central Processing Unit) or MPU (Micro Processing Unit) **as efficiently as possible**. Most kernels are written in C and require a small portion of code written in assembly language in order to adapt the kernel to different CPU architectures.

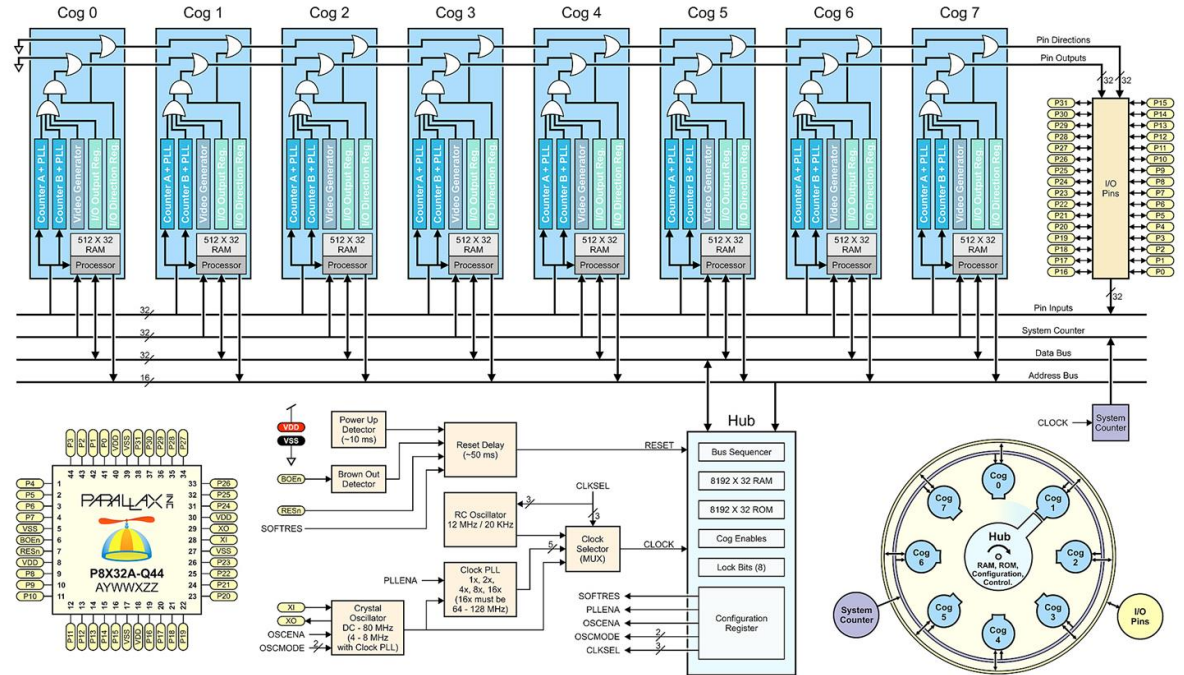

- *void MyTask (void)*

```
{  
    while (1) {  
        Wait for an event to occur;  
        Perform task operation;  
    }  
}
```



Parallax Propeller

- Нет понятия прерываний. Вместо это предлагается запускать конкурирующие задачи разных ядрах (cog'ax).



Hub and Cog Interaction

RealTime and Interrupts

- *Difference between real time operating system and non real time operating system?*



Real time OS

1. A real-time operating system is an operating system intended to serve real-time applications that process data as it comes in, typically without buffer delays.
2. It is deterministic.
3. It is time sensitive.
4. It can't use virtual memory.
5. It is dedicated to single work.
6. It has flat memory model.
7. It has low interrupt latency.

RealTime and Interrupts

Two Real Time Operating System

Hard Real-Time System

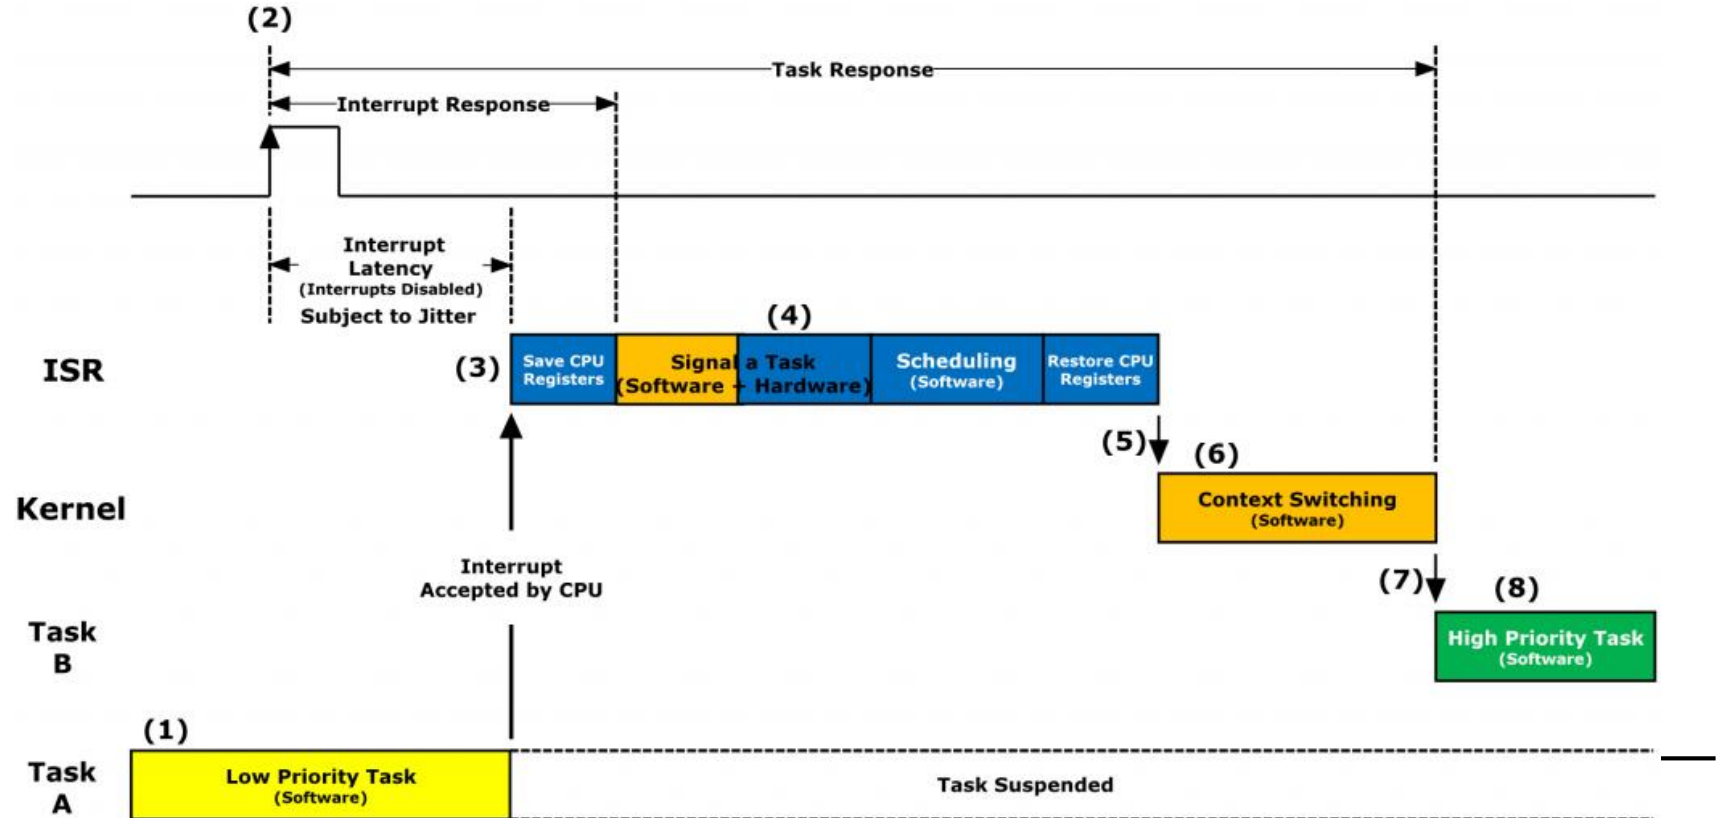
Soft Real-Time System



RealTime and Interrupts

RTOS

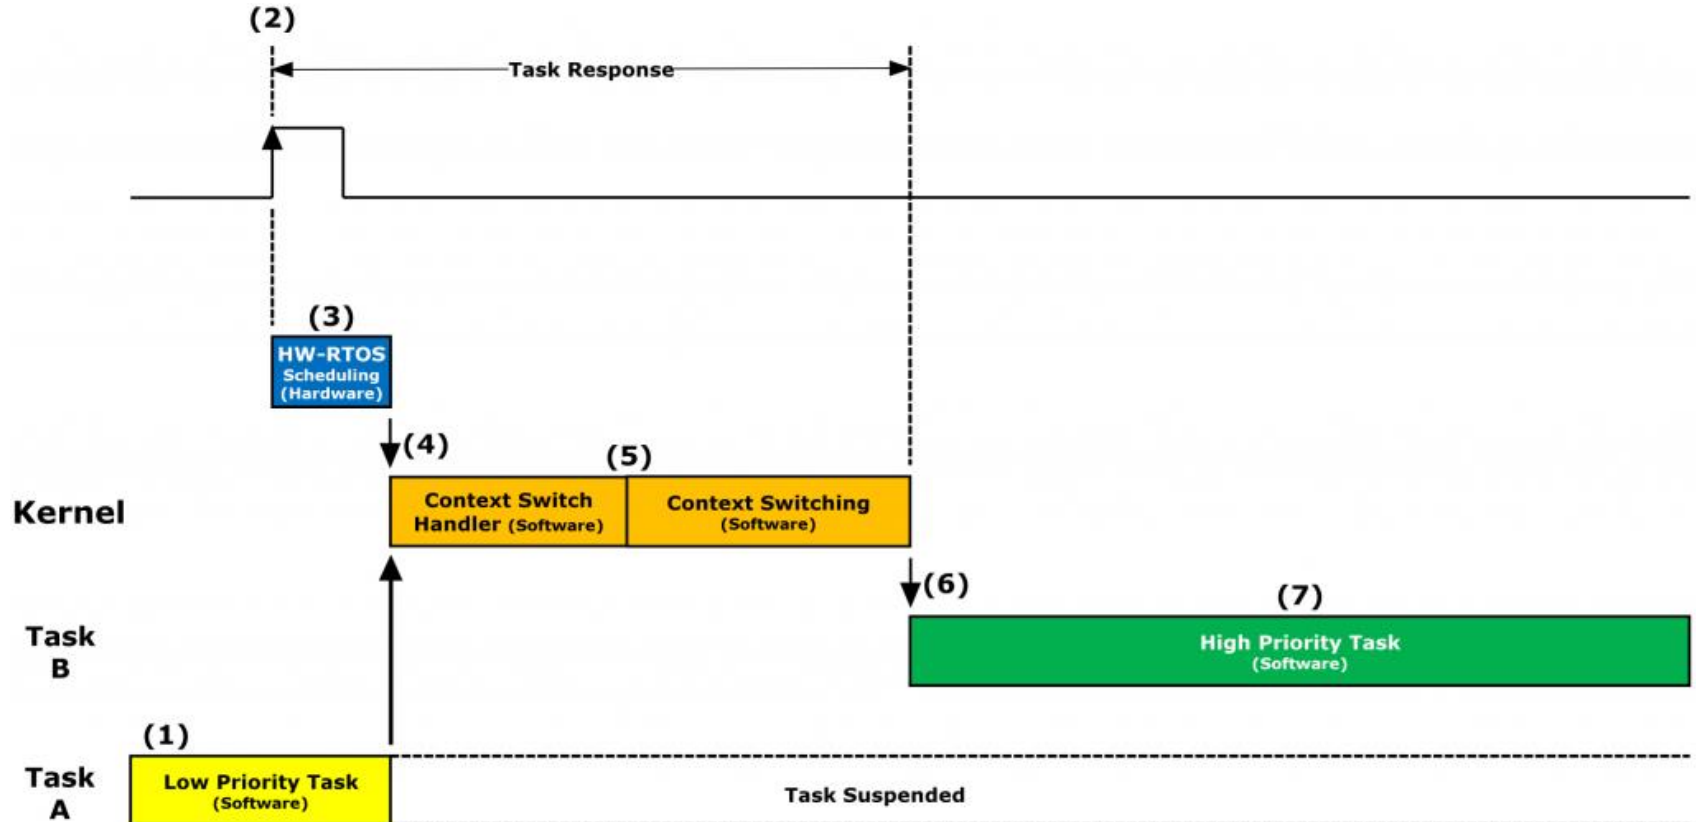
- HW-RTOS can handle interrupting devices using ISRs like other kernels



RealTime and Interrupts

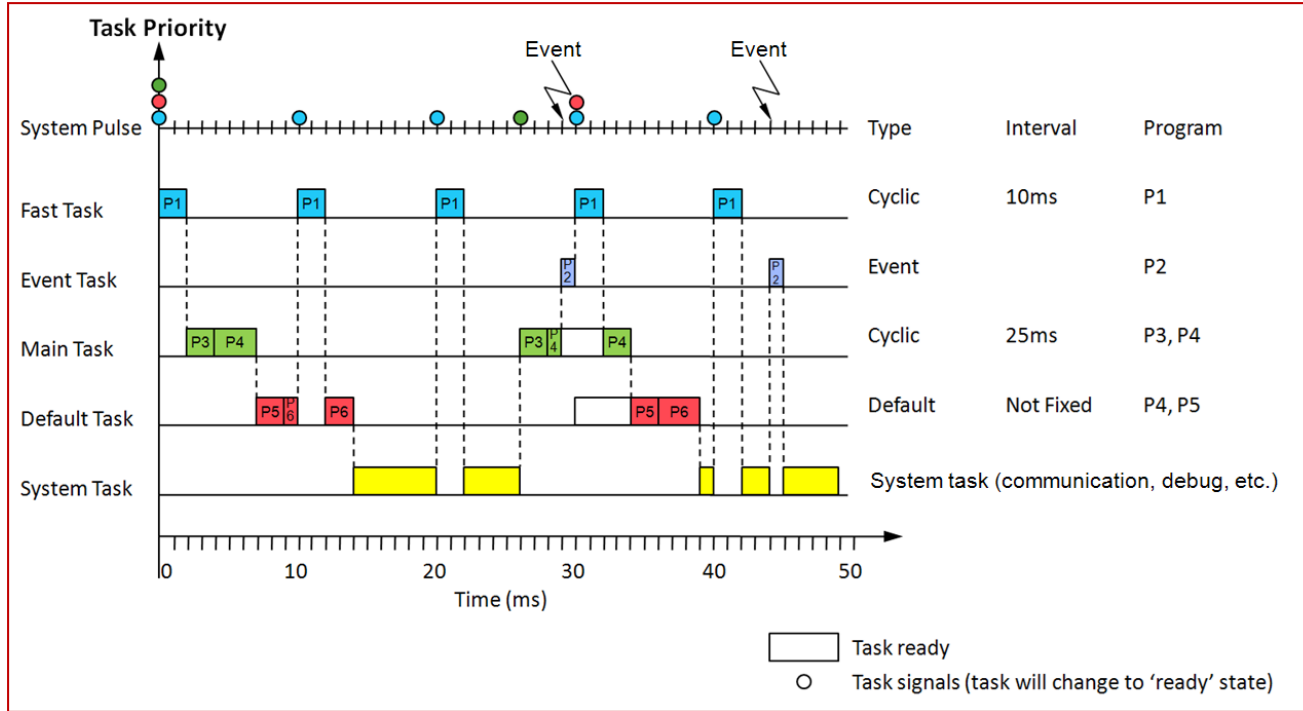
RTOS

- Interrupting devices can directly signal a task



RealTime and Interrupts


- IEC 61131-3 programming 5 tasks in one project.



RealTime and Interrupts

PLC System

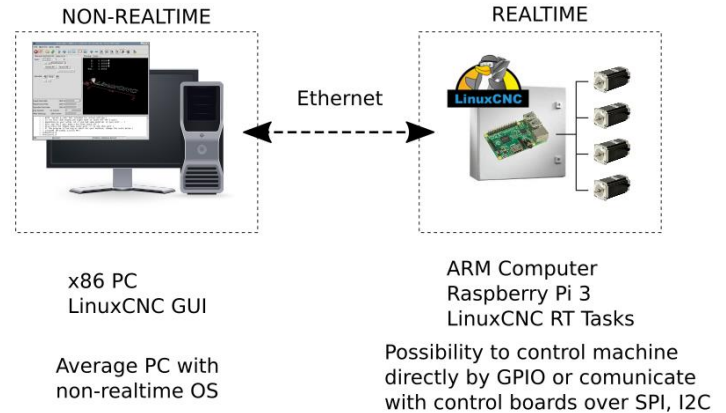
- PLC model



RealTime and Interrupts

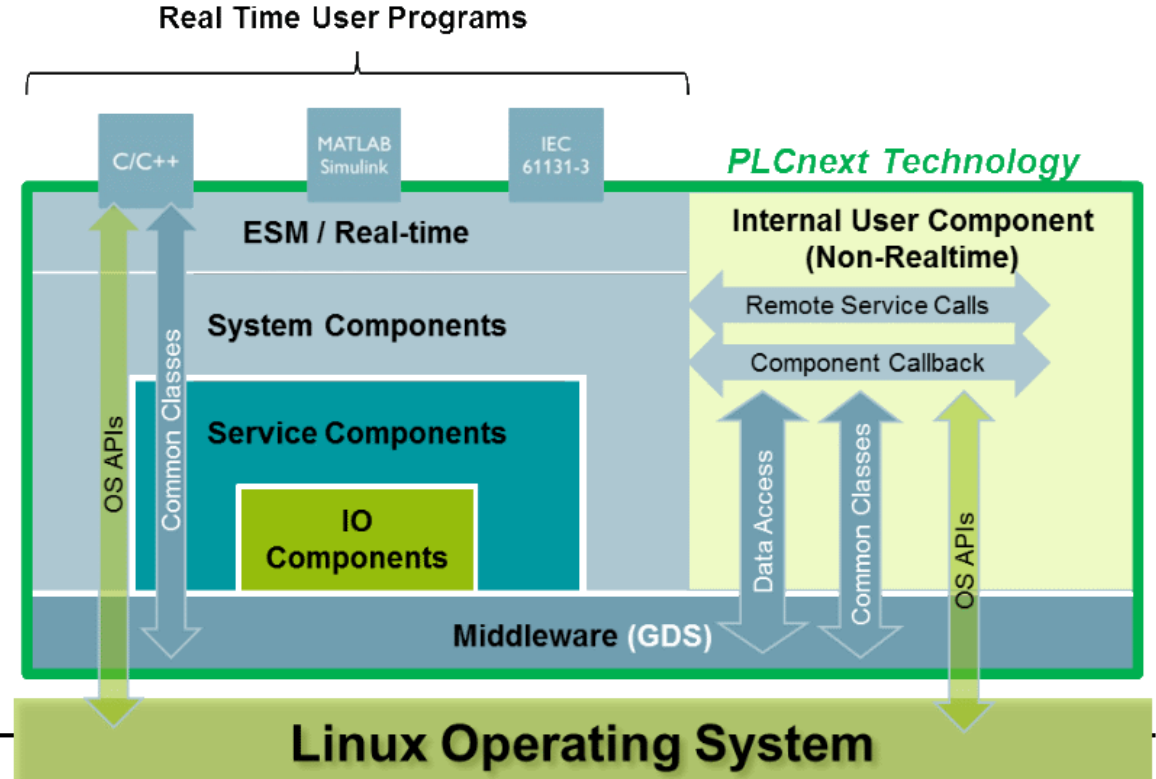
Non-real time OS

1. A Non-real time OS or General purpose OS is the operating system made for high end, general purpose systems like a personal computer, a work station, a server system etc.
2. It is not deterministic.
3. It is time insensitive.
4. It can use virtual memory concept.
5. It is used in multi-user environment.
6. It has protected memory model.
7. It has high interrupt latency




RealTime and Interrupts

- PLCnext model



RealTime and Interrupts

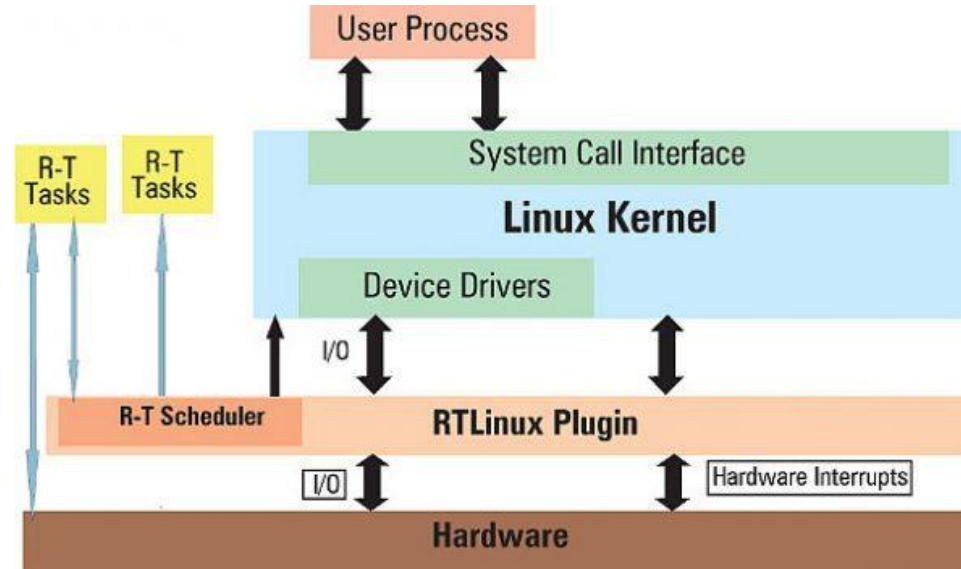
- PLC next



RealTime and Interrupts

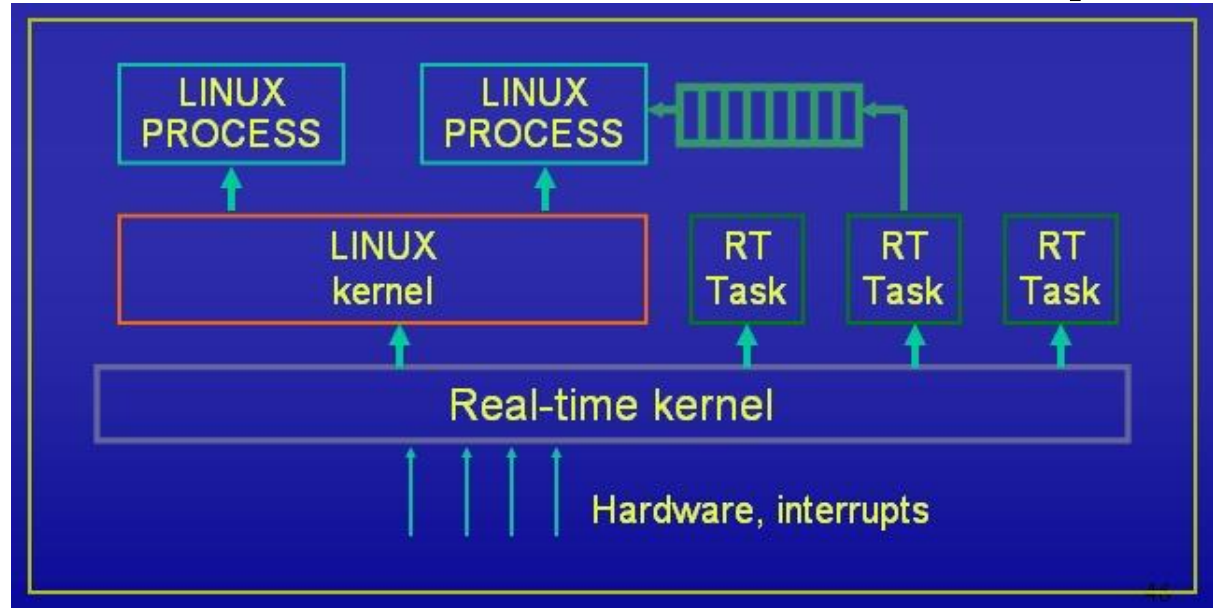
- **RTLinux**

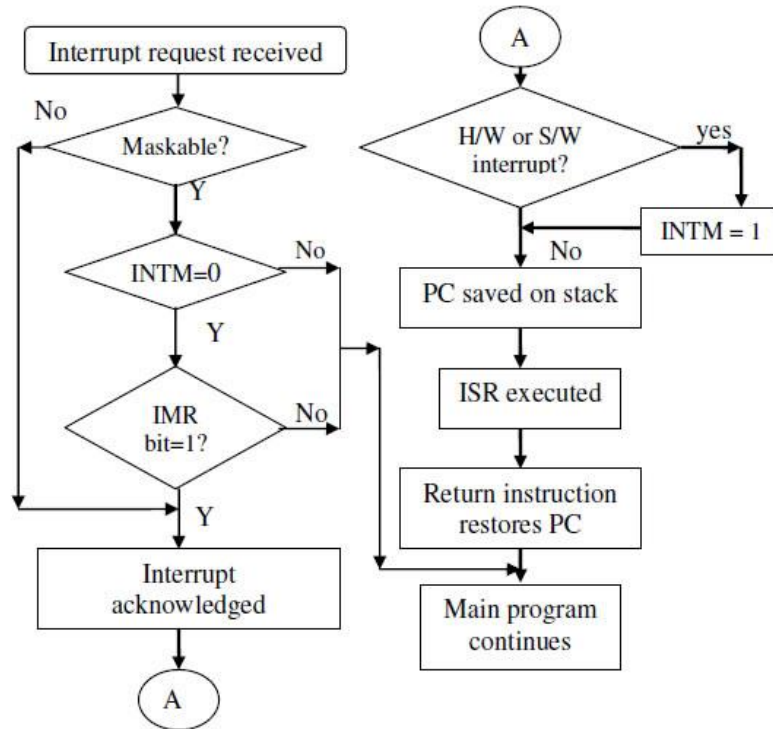
RTLinux is a micronuclear hard real-time operating system that runs Linux as a completely extruded process.



- RTLinux


Hardware interrupt





Maskable interrupts

- **Maskable interrupt (IRQ):** a **hardware interrupt** that may be ignored by setting a bit in an interrupt mask register's (IMR) bit-mask.



Maskable interrupts

- All Interrupt Requests (IRQs) issued by I/O devices give rise to **maskable interrupts**. A maskable interrupt can be in two states: **masked or unmasked**;
 - a masked interrupt is ignored by the control unit as long as it remains masked.
-

IRQ

- Each IRQ line is assigned a numeric value. For example, on the classic PC, IRQ zero is the timer interrupt and IRQ one is the keyboard interrupt.
 - Some interrupts are dynamically assigned, such as interrupts associated with devices on the PCI bus. Other non-PC architectures have similar dynamic assignments for interrupt values.
 - The kernel knows that a specific interrupt is associated with a specific device. The hardware then issues interrupts to get the kernel's attention.
-

IRQ

A problem has been detected and windows has been shut down to prevent damage to your computer.

DRIVER_IRQL_NOT_LESS_OR_EQUAL

If this is the first time you've seen this Stop error screen, restart your computer. If this screen appears again, follow these steps:

Check to make sure any new hardware or software is properly installed. If this is a new installation, ask your hardware or software manufacturer for any windows updates you might need.

If problems continue, disable or remove any newly installed hardware or software. Disable BIOS memory options such as caching or shadowing. If you need to use Safe Mode to remove or disable components, restart your computer, press F8 to select Advanced startup options, and then select Safe Mode.

Technical information:

*** STOP: 0x000000d1 (0xE10C5678, 0x00000002, 0x00000000, 0xF8C48579)

*** myfault.sys - Address F8C48579 base at F8C48000, DateStamp 4d26670c

```
1000100e368 fffff81001b678798 fffff81000100e3b8 fffff81000100e360 7fffffff
1435969.3885561

[41412.679936] [<c10526b0>] ? local_bh_enable+0x90/0x90 000006
[41412.679936] <IRQ>
[41412.679936] [<c10529c5>] ? irq_exit+0x95/0xa0
[41412.679936] [<c161b97b>] ? do_IRQ+0x4b/0xc0
[41412.679936] [<c105298c>] ? irq_exit+0x5c/0xa0
[41412.679936] [<c161ba4e>] ? smp_apic_timer_interrupt+0x5e/0x8d
[41412.679936] [<c161b773>] ? common_interrupt+0x33/0x38
[41412.679936] [<c161007b>] ? export_array+0x3/0x8f
[41412.679936] [<c1610000>] ? i8042_pnp_exit+0x21/0x3e
[41412.679936] Code: 00 8b 53 04 8b 45 d4 89 55 c4 0f b7 40 02 83 c0 07 83 f8 05
04 90 ba 02 00 00 00 <8b> 08 ff 51 34 89 45 e4 8b 03 8b 40 10 83 78 04 01 75 0f
[41412.679936] EIP: [f8b17c55] ulc_dotxstatus+0x6a/0x781 [u1] SS:ESP 0068:f5c0b
[41412.679936] CR2: 0000000000000000
[41412.739322] Kernel panic - not syncing: Fatal exception in interrupt 92 01 0
[41412.743179] drm_kms_helper: panic occurred, switching back to text console 18> 8b 0

er!
er!
```

Non- Maskable interrupts

- Only a few critical events (such as **hardware failures**) give rise to nonmaskable interrupts;
 - Nonmaskable interrupts are **always recognized by the CPU**.
-

Exceptions

- Processor-detected exceptions
- Faults
- Traps
- Aborts
- Programmed exceptions



Exceptions


- Processor-detected exceptions

Generated when the CPU detects an anomalous condition while executing an instruction. These are further divided into three groups, depending on the value of the eip register that is saved on the Kernel Mode stack when the CPU control unit raises the exception.



Exceptions

- **Faults (Ошибки)**
Can generally be corrected; once corrected, the program is allowed to restart with no loss of continuity.
- The saved value of **eip** is the address of the instruction that caused the fault, and hence that instruction can be resumed when the exception handler terminates.



Exceptions

- **Traps (Ловушки)**
- Reported immediately following the execution of the trapping instruction;
- after the kernel returns control to the program, it is allowed to continue its execution with no loss of continuity.
- The saved value of eip is the address of the instruction that should be executed after the one that caused the trap. A trap is triggered only when there is no need to reexecute the instruction that terminated. The main use of traps is for debugging purposes.
- The role of the interrupt signal in this case is to notify the debugger that a specific instruction has been executed. Can generally be corrected; once corrected, the program is allowed to restart with no loss of continuity.
- The saved value of eip is the address of the instruction that caused the fault, and hence that instruction can be resumed when the exception handler terminates.



Exceptions

- **Aborts (Аварии)**
- A serious error occurred;
- The control unit is in trouble, and it may be unable to store in the eip register the precise location of the instruction causing the exception.
- Aborts are used to report severe errors, such as hardware failures and invalid or inconsistent values in system tables.
- The interrupt signal sent by the control unit is an emergency signal used to switch control to the corresponding abort exception handler.
- This handler has no choice but to force the affected process to terminate.

```
An exception occurred

Processor:      ARM11 (core 0)
Exception type: data abort
Current process: menu (0004003000000F02)

R0      00000000      R1      00000000
R2      00000001      R3      39032000
R4      00000000      R5      00000001
R6      384C5290      R7      00320100
R8      38C55B08      R9      38381A30
R10     38C55CCA      R11     00000015
R12     003FFFF0      SP      0FFFFE00
LR      001C5ED8      PC      001C5ED8
CPSR    60000010      FPExc  00000005

You can find a dump in the following file:
dumps/arm11/crash_dump_00000005.dmp

Press any button to shutdown
```


Exceptions

- **Programmed exceptions**
- Occur at the request of the programmer. They are triggered by int or int3 instructions; the into (check for overflow) and bound (check on address bound) instructions also give rise to a programmed exception when the condition they are checking is not true.
- Programmed exceptions are handled by the control unit as traps; they are often called software interrupts .



APIC


APIC - Advanced Programmable Interrupt Controller)



CS	1	28	Vcc
WR	2	27	A0
RD	3	26	INTA
D7	4	25	IR7
D6	5	24	IR6
D5	6	23	IR5
D4	7	22	IR4
D3	8	21	IR3
D2	9	20	IR2
D1	10	19	IR1
D0	11	18	IR0
CAS0	12	17	INT
CAS1	13	16	SP/EN
gnd	14	15	CAS2

APIC was used in multi-core / multi-processor systems, starting with Intel Pentium (core P54). Starting with this processor, each one was supplied with an integrated Local APIC.

PIC

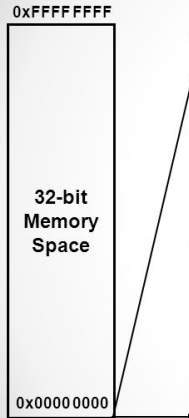


PIC (*Programmable Interrupt Controller*)

APIC was used in multi-core / multi-processor systems, starting with **Intel Pentium (core P54)**. Starting with this processor, each one was supplied with an integrated **Local APIC**.

ARM

The ARM Vector table



0xFFFFFFF

32-bit Memory Space

0x00000000

Address	Vector
0x0000 0000	RESET
0x0000 0004	Undefined Instruction
0x0000 0008	SWI Software Interrupt
0x0000 000C	Abort (Prefetch) Bus Error Inst Fetch
0x0000 0010	Abort (Data) Bus Error Data Fetch
0x0000 0014	Reserved (not used)
0x0000 0018	IRQ Interrupt
0x0000 001C	FIQ Interrupt

AVR

Vector table

16.1. Interrupt Vectors in ATmega328/P

Table 16-1. Reset and Interrupt Vectors in ATmega328/P

Vector No	Program Address ⁽²⁾	Source	Interrupts definition
1	0x0000 ⁽¹⁾	RESET	External Pin, Power-on Reset, Brown-out Reset and Watchdog System Reset
2	0x0002	INT0	External Interrupt Request 0
3	0x0004	INT1	External Interrupt Request 0
4	0x0006	PCINT0	Pin Change Interrupt Request 0
5	0x0008	PCINT1	Pin Change Interrupt Request 1
6	0x000A	PCINT2	Pin Change Interrupt Request 2
7	0x000C	WDT	Watchdog Time-out Interrupt
8	0x000E	TIMER2_COMPA	Timer/Counter2 Compare Match A
9	0x0010	TIMER2_COMPB	Timer/Counter2 Compare Match B
10	0x0012	TIMER2_OVF	Timer/Counter2 Overflow
11	0x0014	TIMER1_CAPT	Timer/Counter1 Capture Event


Interrupt Vectors in Linux

Vector table


Vector range	Use
0-19 (0x0-0x13)	Nonmaskable interrupts and exceptions
20-31 (0x14-0x1f)	Intel-reserved
32-127 (0x20-0x7f)	External interrupts (<i>IRQs</i>)
128 (0x80)	Programmed exception for system calls (see Chapter 10)
129-238 (0x81-0xee)	External interrupts (<i>IRQs</i>)
239 (0xef)	Local APIC timer interrupt (see Chapter 6)
240 (0xf0)	Local APIC thermal interrupt (introduced in the Pentium 4 models)
241-250 (0xf1-0xfa)	Reserved by Linux for future use
251-253 (0xfb-0xfd)	Interprocessor interrupts (see the section "Interprocessor Interrupt Handling" later in this chapter)
254 (0xfe)	Local APIC error interrupt (generated when the local APIC detects an erroneous condition)
255 (0xff)	Local APIC spurious interrupt (generated if the CPU masks an interrupt while the hardware device raises it)

Interrupts

Interrupts control:



signal on
interrupt
line



Interrupts

Interrupts control:

<https://lxr.missinglinkelectronics.com/linux/include/linux/interrupt.h>

- [include/linux/interrupt.h](#)
 - **disable_irq()**
 - **enable_irq()**
 - etc.
- [include/linux/irqflags.h](#)
 - **local_irq_disable()**
 - **local_irq_enable()**
 - **local_irq_save(flags)**

Procs interface:

- /proc/interrupts
 - /proc/irq/
-

Interrupts

- /proc/stat
- /proc/interrupts

```
geeko@dal:~/Documents> cat /proc/interrupts
CPU0
 0:          91   IO-APIC-edge     timer
 1:        1708   IO-APIC-edge     i8042
 3:           1   IO-APIC-edge
 4:           1   IO-APIC-edge
 6:           5   IO-APIC-edge     floppy
 7:           0   IO-APIC-edge     parport0
 8:           0   IO-APIC-edge     rtc0
 9:           0   IO-APIC-fastEOI    acpi
12:        3608   IO-APIC-edge     i8042
14:           0   IO-APIC-edge     ata_piix
15:       61463   IO-APIC-edge     ata_piix
16:           0   IO-APIC-fastEOI    ehci_hcd:usb1
17:       17659   IO-APIC-fastEOI    ioc0
18:        1124   IO-APIC-fastEOI    vmxnet ether
19:           0   IO-APIC-fastEOI    vmci, uhci_hcd:usb2
NMI:           0   Non-maskable interrupts
LOC:       612282   Local timer interrupts
RES:           0   Rescheduling interrupts
CAL:           0   function call interrupts
TLB:           0   TLB shootdowns
TRM:           0   Thermal event interrupts
SPU:           0   Spurious interrupts
```

IRQ interface

Interrupts control:

- [include/linux/interrupt.h](#)
 - **disable_irq()**
 - **enable_irq()**
 - etc.
- [include/linux/irqflags.h](#)
 - **local_irq_disable()**
 - **local_irq_enable()**
 - **local_irq_save(flags)**

Procs interface:

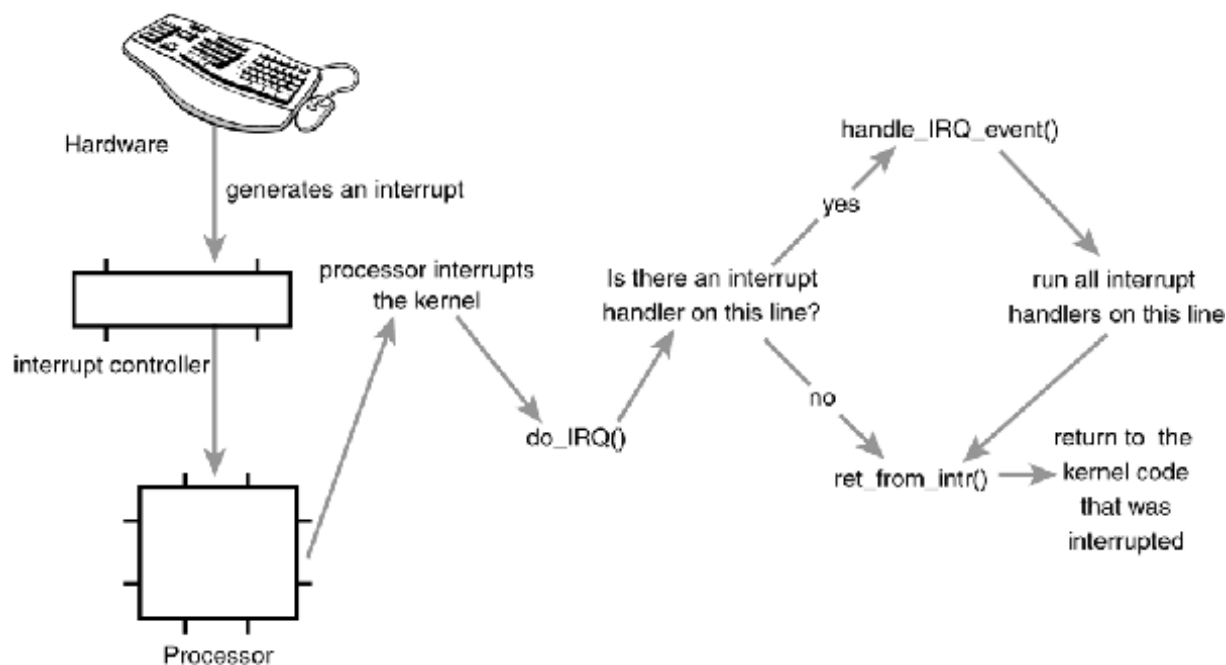
- `/proc/interrupts`
 - `/proc/irq/`
-

IRQ interface

Interrupts control:


- [include/linux/interrupt.h](#)
 - **disable_irq()**
 - **enable_irq()**
 - etc.
 - [include/linux/irqflags.h](#)
 - **local_irq_disable()** - запретить прерывания на локальном CPU;
 - **local_irq_enable()** - разрешить прерывания на локальном CPU
 - **local_irq_save(flags)**
 - **int irqs_disabled()** - вернуть ненулевое значение, если запрещены прерывания на локальном CPU, в противном случае возвращается нуль;
-

IRQ interface



APIC

The diagram illustrates the APIC architecture. At the top, two CPU blocks are shown, labeled CPU₁ and CPU₂. Each CPU block contains a smaller block labeled 'Local APIC'. To the right of the CPU₂ block, an arrow points to its 'Local APIC' block with the label 'Local interrupts'. Below the CPUs is a horizontal double-headed arrow labeled 'ICC bus'. Below the bus is a block labeled 'I/O-APIC'. An arrow points from the 'I/O-APIC' block up to the 'ICC bus'. To the right of the 'I/O-APIC' block, four arrows point towards it with the label 'External interrupts'.



Interrupt handler

- [include/linux/interrupt.h](#)
 - `typedef irqreturn_t (*irq_handler_t)(int, void *)`
 - `int request_irq(unsigned int irq, irq_handler_t handler, unsigned long flags, const char *name, void *dev)`
 - `void free_irq(unsigned int, void *)`
 - **irq** - номер линии запрашиваемого прерывания.
 - **handler** - указатель на функцию-обработчик.
 - **flags** - битовая маска опций (описываемая далее), связанная с управлением прерыванием.
 - **name** - символьная строка, используемая в `/proc/interrupts`, для отображения владельца прерывания.
 - **dev** - указатель на уникальный идентификатор устройства на линии IRQ, для не разделяемых прерываний (например шины ISA) может указываться NULL. Данные по указателю dev требуются для удаления только специфицируемого устройства на разделяемой линии IRQ.
-

Interrupt handler

- Для линии **IRQ** регистрируется функция обработчика «верхней половины» (это та же ISR функция по смыслу, «верхняя половина» обработчика), который выполняется при запрещённых прерываниях локального процессора. Именно этой функции передаётся управление при возникновении аппаратного прерывания.
 - Синтаксически функция-обработчик должна иметь строго описанный функциональный тип **irq_handler_t**, и возвращает управление ядру системы традиционным **return**, с возвращаемым значением **IRQ_NONE** или **IRQ_HANDLED**.
 - При возникновении аппаратного прерывания по линии **IRQ** **функция-обработчик получит управление**. Эта функция выполняется в контексте прерывания — это одно из самых важных ограничений, накладываемых Linux, мы не раз будем возвращаться к нему. Перед своим завершением функция-обработчик регистрирует для последующего выполнения функцию нижней половины обработчика, которая и завершит позже начатую работу по обработке этого прерывания...
-

Interrupt handler

- При возникновении аппаратного прерывания по линии IRQ функция-обработчик получит управление. Эта функция выполняется в контексте прерывания — это одно из самых важных ограничений, накладываемых Linux, мы не раз будем возвращаться к нему. Перед своим завершением функция-обработчик регистрирует для последующего выполнения функцию нижней половины обработчика, которая и завершит позже начатую работу по обработке этого прерывания...
 - В этой точке (после return из обработчика верхней половины) **ядро завершает всё взаимодействие с аппаратурой контроллера прерываний, разрешает последующие прерывания, восстанавливает контроллер командой завершения обработки прерывания** (посылает EOI) и возвращает управление из прерывания (из ядра!) уже именно командой iret. После этого будет восстановлен контекст прерванного процесса (потока).
 - **А вот запланированная выше к выполнению функция нижней половины будет вызвана ядром в некоторый момент позже** (часто это может быть и непосредственно после завершения return из верхней половины, но это непредсказуемо), тогда, когда удобнее будет ядру системы. Принципиально важное отличие функции нижней половины состоит в том, что она выполняется уже при разрешённых прерываниях.
-

free_irq()

- Предупреждение относительно удаления обработчика `free_irq()`
 - Если же в модуле при его завершении (выгрузке) вы не выполните явно `free_irq()`, то почти со 100% вероятностью произойдёт следующее:
 - модуль будет выгружен, но вектор прерывания будет установлен на тот адрес, который перед тем занимала зарегистрированная функция **handler()** ...
 - по истечению некоторого времени эта **область памяти будет переписана ядром** под какие-то иные цели...
 - и первое же произошедшее после этого аппаратное прерывание по этой линии **IRQ** приведёт к немедленному краху всей системы.
-

Softirq

All vectors are registered in **softirq_vec** array.

Kernel uses per-CPU threads **ksoftirqd** to launch deferred interrupts.

- [include/linux/interrupt.h](#)
 - struct **softirq_action**
 - void (***action**)(struct softirq_action *)
 - void **open_softirq**(int nr, void (*action)(struct softirq_action *))
 - void **raise_softirq**(unsigned int nr)

Procs interface:

- /proc/softirqs
-

- Создание нового уровня softirq

1. Определить новый индекс (уровень) отложенного прерывания, вписав (файл <linux/interrupt.h>) свою константу вида **XXX_SOFT_IRQ** в перечисление, где-то, очевидно, на одну позицию выше **TASKLET_SOFTIRQ** (иначе зачем переопределять новый уровень и не использовать тасклет?).
2. Во время инициализации модуля должен быть зарегистрирован (объявлен) обработчик отложенного прерывания с помощью вызова **open_softirq()**, который принимает три параметра: этот индекс отложенного прерывания, функция-обработчик и значение поля **data**
3. Функция-обработчик отложенного прерывания должна в точности соответствовать правильному прототипу:
`void xxx_analyze(unsigned long data);`
 1. Зарегистрированное отложенное прерывание.
 2. Затем, в подходящий (не для вас, для системы) момент времени отложенное прерывание начнёт выполняться

Tasklet

Tasklets are simplified interface to schedule deferred tasks.

It's implemented as softirqs (TASKLET_SOFTIRQ or HI_SOFTIRQ).

- [include/linux/interrupt.h](#)
 - [struct tasklet_struct](#)
 - void **tasklet_init**(struct tasklet_struct *t, void (*func)(unsigned long), unsigned long data)
 - **DECLARE_TASKLET**(name, func, data)
 - **DECLARE_TASKLET_DISABLED**(name, func, data)
 - void **tasklet_schedule**(struct tasklet_struct *t)
 - void **tasklet_hi_schedule**(struct tasklet_struct *t)
 - void **tasklet_disable**(struct tasklet_struct *t)
 - void **tasklet_enable**(struct tasklet_struct *t)
 - void **tasklet_kill**(struct tasklet_struct *t)
-



Workqueue

The workqueue is another concept for handling deferred functions.

Workqueue functions run in the context of a kernel process.

Kernel provides per-cpu threads (**kworke**r) to launch scheduled works.

- [include/linux/workqueue.h](#)
 - [struct work struct](#)
 - [struct delayed work](#)
 - INIT_WORK(_work, _func)
 - DECLARE_WORK(n, f)
 - DECLARE_DELAYED_WORK(n, f)
 - DECLARE_DEFERRABLE_WORK(n, f)
 - [bool schedule work\(struct work struct *work\)](#)
 - [bool schedule delayed work\(struct delayed work *dwork, unsigned long delay\)](#)
 - [void flush scheduled work\(void\)](#)
 - [bool cancel delayed work\(struct delayed work *dwork\)](#)
-

Custom workqueues


You can create your own work queues instead of using global ones:

- [include/linux/workqueue.h](#)
 - `struct workqueue_struct *create_workqueue(name)`
 - `void destroy_workqueue(struct workqueue_struct *)`
 - [`bool queue_work\(struct workqueue_struct *wq, struct work_struct *work\)`](#)
 - [`bool queue_delayed_work\(struct workqueue_struct *wq, struct delayed_work *dwork, unsigned long delay\)`](#)
 - [`bool queue_work_on\(int cpu, struct workqueue_struct *wq, struct work_struct *work\)`](#)
-

Interrupts


gpv@orangepione:/sys/class/gpio\$ cat /proc/interrupts

	CPU0	CPU1	CPU2	CPU3		
17:	0	0	0	0	GICv2 50 Level	/soc/timer@01c20c00
18:	0	0	0	0	GICv2 29 Level	arch_timer
19:	119798	41591	52656	25090	GICv2 30 Level	arch_timer
22:	0	0	0	0	GICv2 120 Level	1ee0000.hdmi, dw-hdmi-cec
23:	0	0	0	0	GICv2 82 Level	1c02000.dma-controller
24:	0	0	0	0	GICv2 118 Level	1c0c000.lcd-controller
25:	250724	0	0	0	GICv2 92 Level	sunxi-mmc
26:	0	0	0	0	GICv2 103 Level	musb-hdrc.1.auto
27:	0	0	0	0	GICv2 104 Level	ehci_hcd:usb1
28:	0	0	0	0	GICv2 105 Level	ohci_hcd:usb2
29:	0	0	0	0	GICv2 106 Level	ehci_hcd:usb3
30:	0	0	0	0	GICv2 107 Level	ohci_hcd:usb4
34:	23984	0	0	0	GICv2 114 Level	eth0
37:	2034	0	0	0	GICv2 32 Level	ttyS0
38:	0	0	0	0	GICv2 72 Level	1f00000.rtc
92:	1	0	0	0	sunxi_pio_edge 44 Edge	usb0-id-det
IPI0:	0	0	0	0	CPU wakeup interrupts	
IPI1:	0	0	0	0	Timer broadcast interrupts	
IPI2:	10088	34948	15124	13968	Rescheduling interrupts	
IPI3:	13	13	11	11	Function call interrupts	
IPI4:	0	0	0	0	CPU stop interrupts	
IPI5:	5030	1574	2250	311	IRQ work interrupts	
IPI6:	0	0	0	0	completion interrupts	
Err:	0					



Interrupts

- `cat /sys/kernel/debug/gpio`
- `gpiochip0: GPIOs 0-223, parent: platform/1c20800.pinctrl, 1c20800.pinctrl:`
- `gpio-13 (|sysfs) out hi`
- `gpio-15 (|orangepi:red:status) out lo`
- `gpio-16 (|sysfs) out lo`
- `gpio-166 (|cd) in lo`
- `gpio-204 (|usb0_id_det) in hi IRQ`
- `gpiochip1: GPIOs 352-383, parent: platform/1f02c00.pinctrl, 1f02c00.pinctrl:`
- `gpio-354 (|usb0-vbus) out lo`
- `gpio-358 (|?) out lo`
- `gpio-362 (|orangepi:green:pwr) out hi`
- `root@orangeone:/sys#`



Do It

- <http://blablacode.ru/yadro-linux/540>
- <http://blablacode.ru/yadro-linux/543>
- https://github.com/lamazavr/linux_kernel_mod
- https://github.com/lamazavr/linux_kernel_mod/tree/master/blablamod_interrupts
- See https://www.youtube.com/watch?v=ltCZydX_zmk&t=775s
- See <https://drive.google.com/drive/folders/1wgZfagxHicM8iz1yVCGtRF5sQxXpKZlw>



Do It

- Создание "нижней половины" обработчика прерываний
https://www.ibm.com/developerworks/ru/library/l-linux_kernel_59/index.html
- Создание "верхней половины" обработчика прерываний
https://www.ibm.com/developerworks/ru/library/l-linux_kernel_58/
- Тасклеты и очереди отложенных действий
https://www.ibm.com/developerworks/ru/library/l-linux_kernel_60/

